

Natural Language and Speech Processing

Lecture 12: Tips and Tricks

Tanel Alumäe

Contents

- Tips about data
- Tips about models
- Tips about training
- Debugging neural networks
- Other practical tips

Partly based on Andrej Karpathy's (Head of AI at Tesla) recent blog post <http://karpathy.github.io/2019/04/25/recipe/> (must read)

Data

- Become one with data
 - Inspect your data carefully (spend many hours)
 - Understand the distribution and look at the patterns
 - Your brain is very good at it
 - Are there duplicates?
 - Are there corrupt examples?
 - Are the labels imbalanced?
 - Think how your brain solves this task
 - Can you solve this task on local context only?
 - Or do you need larger (global) context?

How much data?

- Very general rule: you can start thinking about training a DNN if you have at least 1000 training samples
- **There is no data like more data**
- Common situation:
 - Baseline accuracy with a simple model: accuracy 60%
 - Super-well tuned “transformer with multi-head attention”: accuracy: 70%
 - Spend 1 week
 - Simple model with more data: 80%
 - Spend 2 days on data collection
- Spending time on data collection is a time well spent
 - Can always apply new models / training tricks when they are invented
 - Time spent on fine-tuning a fancy model might be wasted when suddenly a new model architecture is invented

Data augmentation

- *“The next best thing to real data is half-fake data”*
- Data augmentation: train artificial additional training data by corrupting the real data by some domain-specific way
- Often applied so that the amount of training data is increased 10 (or more) times
- Data augmentation can make a huge difference, especially if domain adaptation can be performed using augmentation

Data augmentation, continued

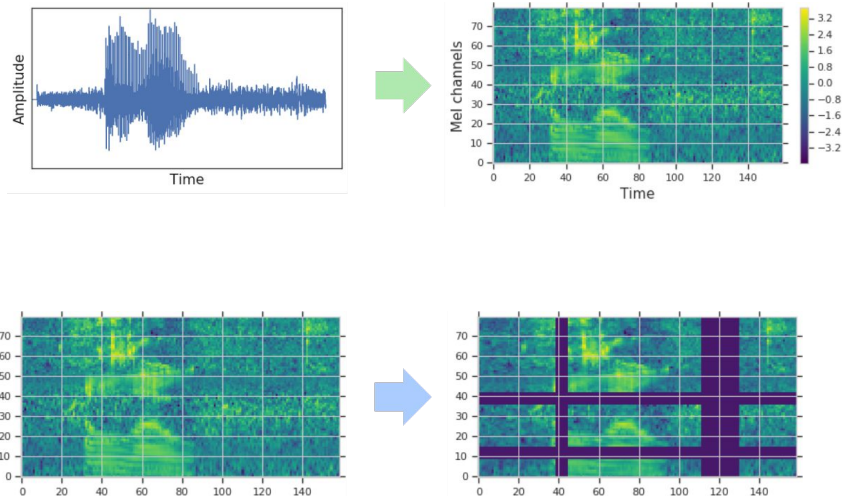
- For example:
 - Speech recognition model trained on clean data: accuracy on reverberant/noisy data **50%**
 - Add noise and reverberate clean training data: accuracy **85%**
- Data augmentation is more useful in domains where input data is “continuous”
 - Image processing
 - Speech processing
- Usually works better than trying to “fix” the test data (e.g., denoise)
- Almost never hurts accuracy

Data augmentation is speech recognition

- Speed perturbation
 - Add copies of training sentences, sped up or slowed down by 10%
 - Simulates faster/slower speaking, but also different vocal tract characteristics
- Reverberation
 - Add room effects (using real or artificial **impulse responses**)
- Noise augmentation
 - Add background noise to (clean training data), e.g. street noise, engine, car, wind
 - Add “foreground noises” (e.g., door slams, claps, clicks)
 - Add music
 - Add “babble noise” (many persons speaking in the background)
 - All this with random signal-to-noise ratio, and many times

Spectral augmentation


- Proposed by Google recently (<https://ai.googleblog.com/2019/04/specaugment-new-data-augmentation.html>)
- Three types of modifications of the raw filterbank spectrogram, randomly chosen
 - Time warping
 - Masking blocks of consecutive filterbank features (horizontal)
 - Masking whole filterbank blocks in time (vertical)



	LibriSpeech 960h		Switchboard 300h	
	test-clean	test-other	Switchboard	CallHome
Previous SOTA	2.95	7.50	8.3	17.3
Our Results	2.5	5.8	6.8	14.1

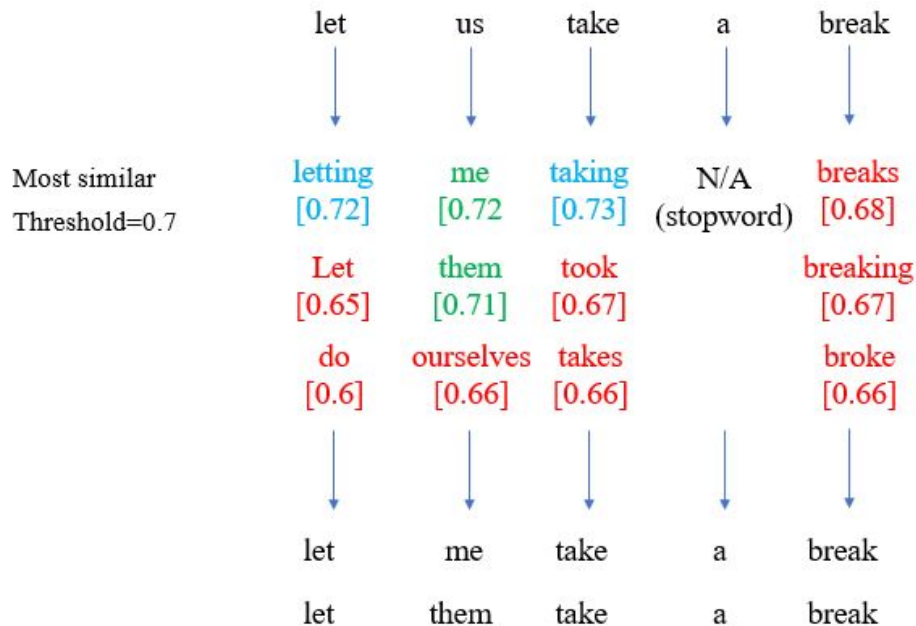
Data augmentation with images

- Flipping
- Rotation
- Translation
- Scaling
- Adding noise
- ...



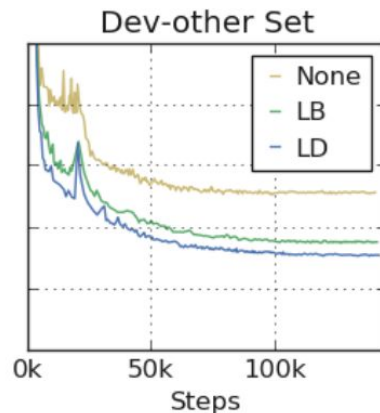
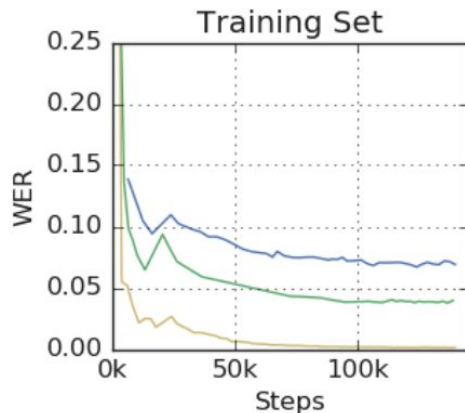
Data augmentation for NLP

- For text data, data augmentation is less useful
- One idea: replace words in training data with words that have similar embeddings
 - But usually better to just use pretrained word embeddings (or pretrained contextualized representations, like BERT)
- Alternative: use domain knowledge
 - E.g., replace company names in text with other company names from the business registry (if having a good coverage of company names is relevant to your task)



Data augmentation under-fits training data

- Data augmentation turns over-fitting problem into under-fitting problem
 - Loss when training augmented data is worse than when training on clean data
 - But loss (and error rate) on test data is better
- Thus, common methods that work against under-fitting help:
 - Using larger (wider and deeper) models
 - Training for more epochs




Mixup


- During training, mix two training samples and the corresponding labels

$$\tilde{x} = \lambda x_i + (1 - \lambda)x_j, \quad \text{where } x_i, x_j \text{ are raw input vectors}$$
$$\tilde{y} = \lambda y_i + (1 - \lambda)y_j, \quad \text{where } y_i, y_j \text{ are one-hot label encodings}$$

- The mixup factor is drawn randomly from a beta distribution



- Kind of surprising that it helps
- Not really applicable to text data



Pretraining

- *It rarely ever hurts to use a pretrained network if you can, even if you have enough data*
- Pretraining was first popularized by the image recognition community
 - Pretrain a ConvNet on ImageNet
 - ImageNet: dataset of 1.2M images, 1000 categories
 - Models trained on ImageNet can be used to initialize models **for completely other datasets** and improve performance significantly
 - Works even if you have only a few examples per object

Image classification

Easiest classes

red fox (100) hen-of-the-woods (100) ibex (100) goldfinch (100) flat-coated retriever (100)



tiger (100)




hamster (100)



porcupine (100)



stingray (100)



Blenheim spaniel (100)



Hardest classes

muzzle (71)



hook (66)



hatchet (68)



spotlight (66)



water bottle (68)



ladle (65)



velvet (68)



restaurant (64)




loupe (66)



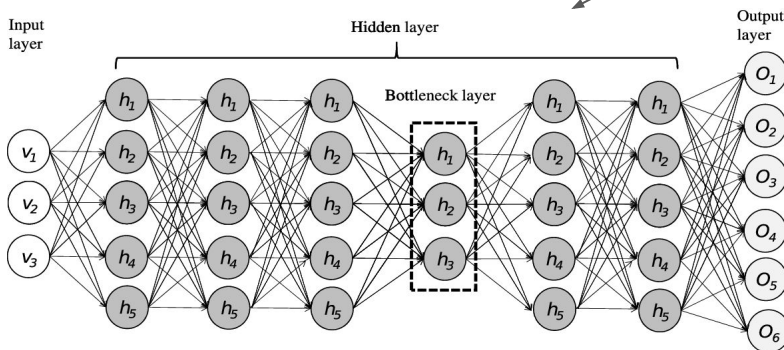
Pretraining for NLP

- Pretraining can have huge benefits on NLP tasks
 - Pretrained embeddings
 - Pretrained contextualized embeddings (e.g. BERT)
- Use pretrained model “found” from internet (many Estonian models available) or train your own
- Cross-lingual pretraining
- Main idea:
 - Train word embeddings for two languages (e.g., English and Italian)
 - Align English and Italian embeddings
 - Use English labeled data to train a text classification model
 - The model also works for Italian!
 - <https://github.com/facebookresearch/MUSE>



Pretraining for speech recognition

- Pretraining also works for speech recognition
- Example:
 - Train a model on a large *universal* corpus
 - Use a small corpus (e.g. 20 h) of in-domain data to finetune the model



- Alternative: (multilingual) **bottleneck** features
- Useful in a situation when you have very little data (e.g. 40h) from a particular language
- Idea:
 - Train a DNN acoustic model on other language(s)
 - The DNN has a narrow *bottleneck* layer before the output
 - Use the bottleneck as a **feature extractor**
 - Train a speech recognition model on the low-resource language on bottleneck features, instead of the usual filterbanks

Pretraining for speaker recognition

- Pretraining is **crucial** for speaker recognition (e.g., speaker identification and verification)
- Idea:
 - Use a large speaker database (speech-speaker pairs) to train a speaker identification model (convolutional model with global pooling)
 - Use the model's pooling layer output as a feature extractor
 - Now, use the feature extractor to get *speaker embeddings* (called x-vectors) for speakers you need to cover
 - Only a few sentences per speaker is now enough

VoxCeleb

- But how to get the “large speaker database”, needed for training the voice feature extractor?
- VoxCeleb:
 - Query YouTube, using e.g. “Elon Musk interview”
 - Use Face Recognition to identify frames where Elon Musk is in the frame
 - Use lip-syncing to check that lips and audio are in sync
 - Extract all such frames and use this as training data
 - Alternative: given several videos of Elon Musk, the segments where he is speaking can be found automatically (invented in our lab, based on a Master thesis)



What model to use in NLP?

- What model to use for NLP task (e.g. text classification, information extraction), given that you have reasonable amount of training data?
- For most NLP tasks, **Transformer** (multi-layer multi-head attention with position encodings) is currently the best choice
- For English, use BERT for getting contextualized word embeddings
- For Estonian, train your own BERT?



Aran Komatsuzaki

@arankomatsuzaki

Follow



Whenever I hear the name 'LSTM,' I can't help but wonder if the person is still living in pre-2017.

6:13 AM - 28 Apr 2019

1 Like



Debugging neural networks

- *Neural net training fails silently*
- Tricky to unit test
- Your net can still (shockingly) work pretty well even if your training data is corrupted somehow
 - E.g., maybe you feed it columns instead of rows of data
 - It's because neural nets can memorize the data to some extent
- Most of the time when you screw something up, it will train but silently work a bit worse

Recipe for training neural nets

- Verify your data
- Start with simple models
 - No data augmentation, fancy learning rate schedules, etc
- Train an input-independent baseline, (e.g. easiest is to just set all your inputs to zero)
 - This should perform worse than when you actually plug in your data without zeroing it out.
Does it?
- Overfit a single batch of only a few examples (e.g. as little as two).
 - Verify that we can reach the lowest achievable loss (e.g. zero)

Recipe, continued



- Don't be a hero
 - Start with a model that is known to give good results on similar tasks
 - You're allowed to do something more custom later and beat this
- One complexity at a time
 - Add fanciness to your model one by one and make sure it improves the model
- Leave learning rate tuning to the end

Recipe, continued

- Regularize (avoid overfitting)
 - Try to get more data
 - Data augmentation, try to be creative
 - Pretraining rarely hurts
 - Unsupervised learning usually gives very little improvement, and is very complicated
 - Try with less features
 - Try smaller model size
 - Add dropout
 - Add L2 regularization (weight decay)
 - Early stopping
 - Try larger model, but stop early

Some tricks: learning rate schedule

- Learning rate schedule
 - E.g., step decay (0.1 at the beginning, 0.0001 at the end)
 - Or: constant learning rate, but decrease it when the loss on dev data doesn't decrease any more (ReduceLROnPlateau in Pytorch)
 - More recent: triangular learning rates
 - Start at learning rate 0.001
 - Gradually increase it to 0.01 at epoch 25%
 - Then gradually decrease it until the end
 - Cyclic learning rates



Some tricks: dropout, label smoothing

- Similarly to learning rate, dropout can be also scheduled
 - E.g.: 0 dropout at the beginning,
 - Increase dropout gradually to 0.3 at epoch 50%
 - Then gradually decrease it back to 0
- Label smoothing
 - Very simple idea: redistribute 0.1 of probability mass from the correct label to all other labels
 - Often used in Google's papers
 - Results in less loss fluctuations on dev data during epochs in my experience
 - Maybe disable label smoothing during the last training phase

Ensembles

- *Model ensembles are a pretty much guaranteed way to gain 2% of accuracy on anything*
 - Ensembles of models of different architecture result in the best performance
 - Even if a particular model is relatively bad (compared to the best ones in the ensemble), combining its predictions with others can give surprising gains
- *If you can't afford the computation at test time, try distilling your ensemble into a network using dark knowledge*
 - Compute posteriors (target probabilities) for your training data using the ensemble, and train a final model on the posteriors (instead of real targets)
 - This way, the final model learns to mimic the ensemble